

Chapter 3

Simon Ward-Jones
simonwardjones16@gmail.com

August 24, 2020

[Link to chapter](#)

1 Linear Regression

- Linear model
- Mini batch stochastic gradient descent
- Vectorisation of algorithms leads to massive efficiency gains
- Minimizing the mean squared error is equivalent to maximum likelihood estimation of a linear model under the assumption of additive Gaussian noise

References

Zhang, A., Lipton, Z. C., Li, M., & Smola, A. J. (2020). *Dive into deep learning*. (<https://d2l.ai>)