

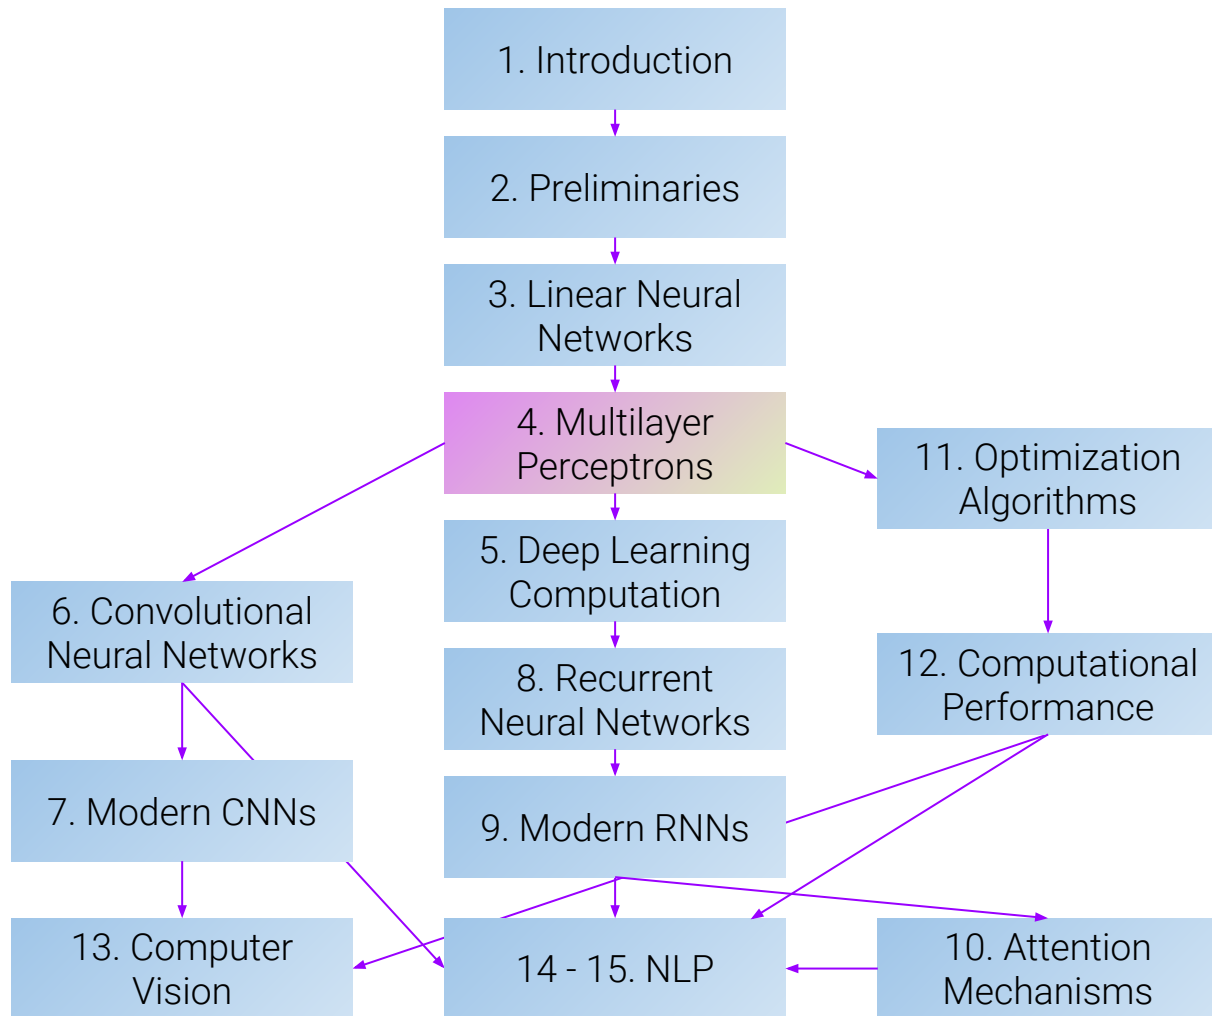


# Multilayer Perceptrons

**Session #4**

*A study group by dair.ai*

# Chapter 3



# Agenda

- Multilayer perceptrons
- Implementation of multilayer perceptrons from scratch
- Model selection, underfitting, overfitting
- Regularization
  - Weight decay
  - Dropout
- Forward and Backward propagation
- Numerical stability and Initialization
- Environment and Distribution Shift
- Predicting house price on Kaggle



# Linearity Assumption


- Linearity implies weaker assumption of ***monotonicity***
  - Increase in features cause
    - Increase in model's output (if weight is positive)
    - Decrease in model's output (if weight is negative)
- E.g., Increasing pixel intensity helps to distinguish between cats and dogs images
  - Inverted images may fail completely
  - Need more robust representations that consider context and relevant feature interactions
- We need a way to model more complex feature relationships

**jointly learn a representation (via hidden layer) and a linear predictor acting on that representation**




# Multilayer Perceptron (MLP)

- Simple deep neural networks are called multilayer perceptrons (MLPs)
- **Multiple layers** of neurons fully connected
  - Hidden layers that learn representations
- Train high-capacity models that can help with **overfitting**




# MLP Architecture

We can overcome limitations of linear models and handle more general functions by adding **hidden layers** in an MLP architecture




2-layer MLP (fully connected) -  
(Figure reproduced from d2l.ai)

# MLP Architecture (Vanilla)



# MLP Architecture (with Activation Function)





# Activation functions

- Decide what neurons will be activated (control the flow of information)
- They are differentiable operators
- Adds **non-linearity** and allows the model to learn complex functions to fit data
- A few activation functions:
  - **ReLU** - produced derivatives that are well behaved and help mitigate *vanishing gradients*
  - **Sigmoid** - squashing function  $\rightarrow (0, 1)$ ; used in binary classification problems
- Proper selection of activation functions become important for **training stable and effective models** (upcoming topic)



# ML Generalization

- Memorization vs. generalization
- When training models on finite datasets there is risk of **memorization**
  - The discovered associations appear to hold but they don't on unseen examples
  - Fail to fit the underlying distribution
  - Leads to **overfitting**
- We desire to train a **generalizable model** that discovers a general pattern
  - Useful for making predictions on future unseen samples
  - Achieved by employing **regularization** techniques that tackle overfitting
- The **training error** is the error of the model on the training dataset
- The **generalization error** is expectation of the model's error
  - We estimate this error on independent **test set** from same underlying data distribution



# Factors that affect model generalizability

- **Tunable parameters**
  - Large number (too many features) could lead to overfitting (model complexity)
- **Parameter values**
  - Wider range of values can lead to overfitting (model complexity)
- **Training examples**
  - More data usually helps in tackling overfitting especially for deep neural networks



# Model selection

- Models can differ in architecture, family, hyperparameters, etc.
- To avoid overfitting on test data, we use the **validation dataset** to determine hyperparameters and eventually select best model
- Compare different models on validation dataset
  - We are interested in generalization error
- In production, this gets more difficult
- When data is scarce, **K-fold cross validation** is employed



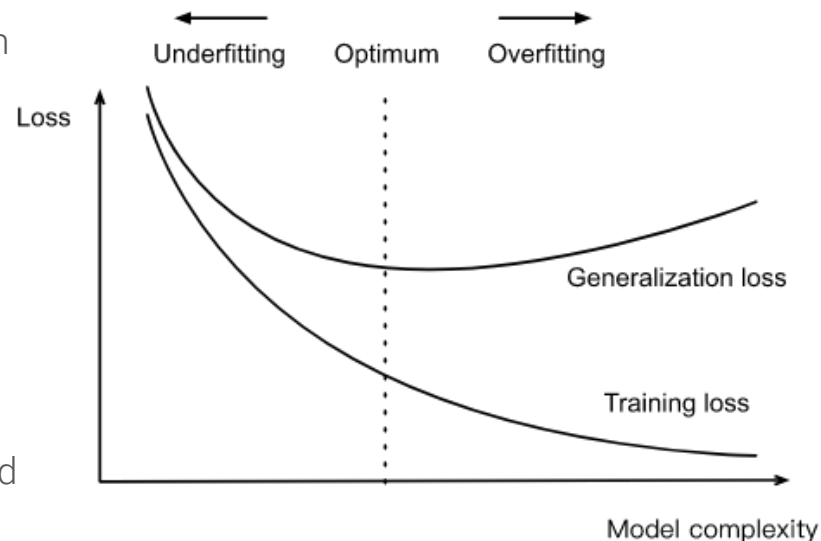
# Underfitting and Overfitting

- **Underfitting**

- Occurs when training and validation error are substantial and there is a small gap between them
- Occurs when working with a simple model
- A more complex model can help

- **Overfitting**

- happens when training error is significantly lower than validation error
- Regularization techniques help
- Increasing # of examples
- Model too complex (flexible and high capacity) and not enough data



# Datasets and deep learning

- When working with small datasets, simpler models can perform better
- Deep learning thrives on large datasets
- However, large datasets can be expensive to collect or time consuming
- Assume we have good quality dataset, we can rely on regularization techniques to avoid overfitting:
  - Weight decay
  - Dropout



# Weight Decay

- We can always reduce overfitting of the model by collecting more data
  - ... but sometimes that process is too expensive or time consuming
- We can also address overfitting through a regularization technique known as weight decay
  - Commonly called L2 regularization and widely used in parametric ML models
  - Measure complexity of linear function
    - We want to ensure that the weight vector doesn't grow too large
    - In other words, we would like to keep a small weight vector (smaller weight values)
  - We add a norm of the weight vector as a penalty term to the original loss function and minimize both
  - Larger weight vector means that the learning algorithm should focus on minimizing the weight norm

$$L(\mathbf{w}, b) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



# L2 vs. L1 Norm

- L1 norm emphasizes on models that concentrate weights on a small set of features while minimizing the effect of others (feature selection)
- L2 norm is computationally convenient and easy to obtain it's derivatives
- L2 norm emphasizes that the model distribute weight evenly across large number of features
  - By penalizing large components of the weight vector
  - Makes model more robust to measurement error in single variable?
- Weight update also considers shrinking of size of weight vector towards zero:

$$\mathbf{w} \leftarrow (1 - \eta\lambda) \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} \left( \mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right)$$





# Dropout

- So far we used techniques of regularization that measure simplicity and reduce model complexity:
  - Tweaking degree of fitted polynomial (limiting number of features)
  - Weight decay (L2 regularization)
- Simplicity can also be represented through **smoothness**
  - Neurons become less sensitive to the activation of another specific neuron
- Dropout enforces smoothness by injecting noise into each layer
  - Drops out some neurons during training by zeroing out fraction of nodes
  - Done on forward propagation and backward propagation
- Dropout break co-adaptation that happens between layers
  - Overfitting happens in a state where each layer depends on specific pattern of activation from previous layers (co-adaptation) → leads to overfitting



# Dropout

$$h' = \begin{cases} 0 & \text{with probability } p \\ \frac{h}{1-p} & \text{otherwise} \end{cases}$$


- Dropout applied by injecting noise through debiasing of each layer
- Debiasing of layer achieved by normalizing using fraction of **nodes retained**
- Apply dropout by zeroing out hidden unit with probability **p**
- Intermediate activation replaced by random variable, **h'**, as follows

```
def dropout_layer(X, dropout):  
    assert 0 <= dropout <= 1  
    # In this case, all elements are dropped out  
    if dropout == 1:  
        return torch.zeros_like(X)  
    # In this case, all elements are kept  
    if dropout == 0:  
        return X  
    mask = (torch.Tensor(X.shape).uniform_(0, 1) >  
            dropout).float()  
    print(mask)  
    return mask * X / (1.0 - dropout)
```

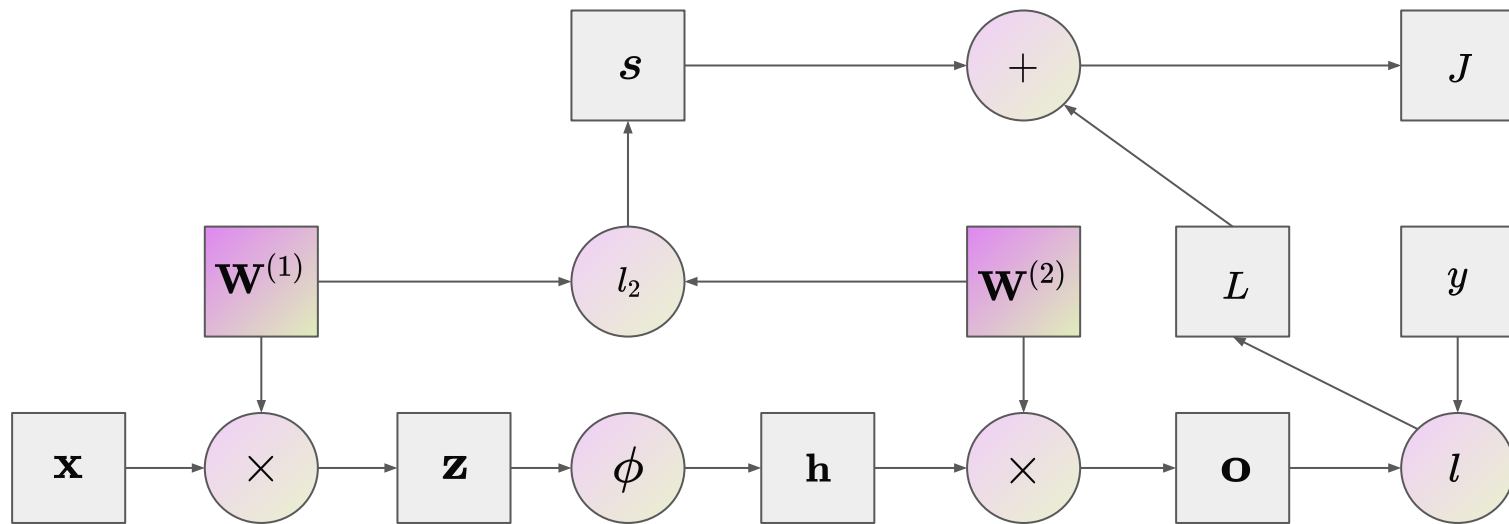


# Dropout intuition

- Reduces the over dependence of output on elements in the hidden layer
- Note that their respective gradients also vanish during backpropagation



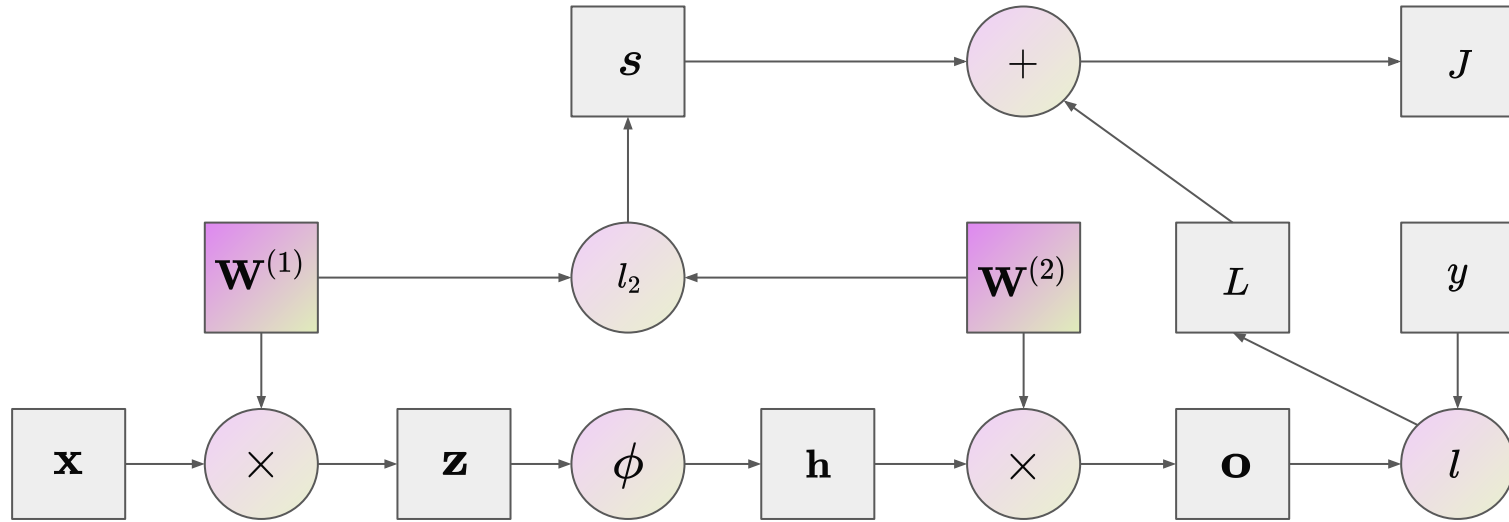
# Forward Propagation



Forward propagation (Figure reproduced from [d2l.ai](https://d2l.ai))



# Backward propagation (Assignment 3)



# Numerical stability and Initialization

- **Initialization plays a huge role in training:**
  - Maintaining numerical stability and ensure parameters & gradients remain well controlled
  - Different initialization schemes can be combined with activation functions in interesting ways
  - Functions and initialization impact how fast optimization algorithm converges
  - If not done right, it may lead to vanishing and exploding gradients
- **Unstable gradients lead to:**
  - Unstable optimization algorithm
  - Parameter updates that are too large... and destroy the model
  - Parameter updates that are too small... and render learning impossible



# Environment and Distribution Shift

- As ML practitioners, it's important to understand ***origin of data*** and what we plan to do with ***model outputs***
- Distribution shift in data can lead to failed models in production
- Users of your ML models may obtain insights/knowledge into how ML model produce outputs and use it to their advantage (e.g., loan repayment criteria)
- Some concerns:
  - Simple models
  - Models that are too technically difficult
  - How about ethical use of algorithms?



# Type of Distribution shift

- **Covariate shift**

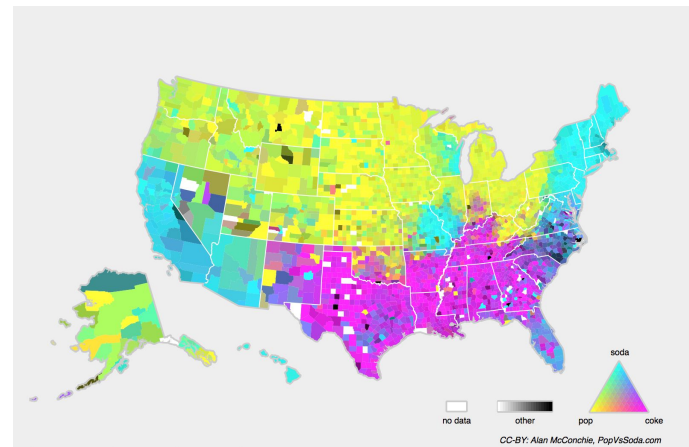
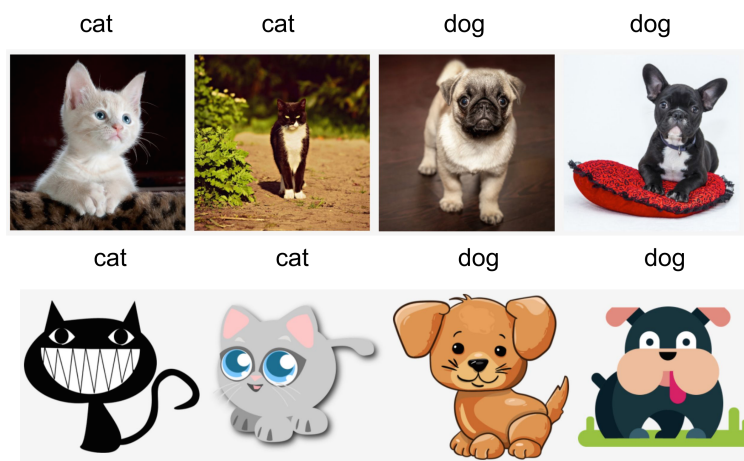
- A shift in the distribution of the features
- Training on realistic photos, testing on cartoons

- **Label shift**

- A shift in the distribution of labels
- Predicting diagnoses based on symptoms
  - Disease cause symptoms

- **Concept shift**

- Definitions of labels change (jobs, mental illness, soft drink name in the US)
- Typically occurs temporally or geographically





# A taxonomy of learning problems

- **Batch learning**
  - Train once, deploy and never update again (app: only let cats into the house)
- **Online learning**
  - Continuously improve learning algorithm based on observations (app: stock price prediction)
- **Bandits**
  - Limit number of actions that can be taken
- **Control**
  - Use automatic control theory (e.g., improve diversity of generated text)
- **Reinforcement learning**
  - Support better interaction with robust environment (e.g., Go, Chess, self-driving car)

# Announcements

- I will post a recording of chapter 5 in the coming days
- Next week we will have Samil delivering CNNs and modern CNNs



Salim Chemlal, Ph.D.  
Deep Learning Instructor at NVIDIA