

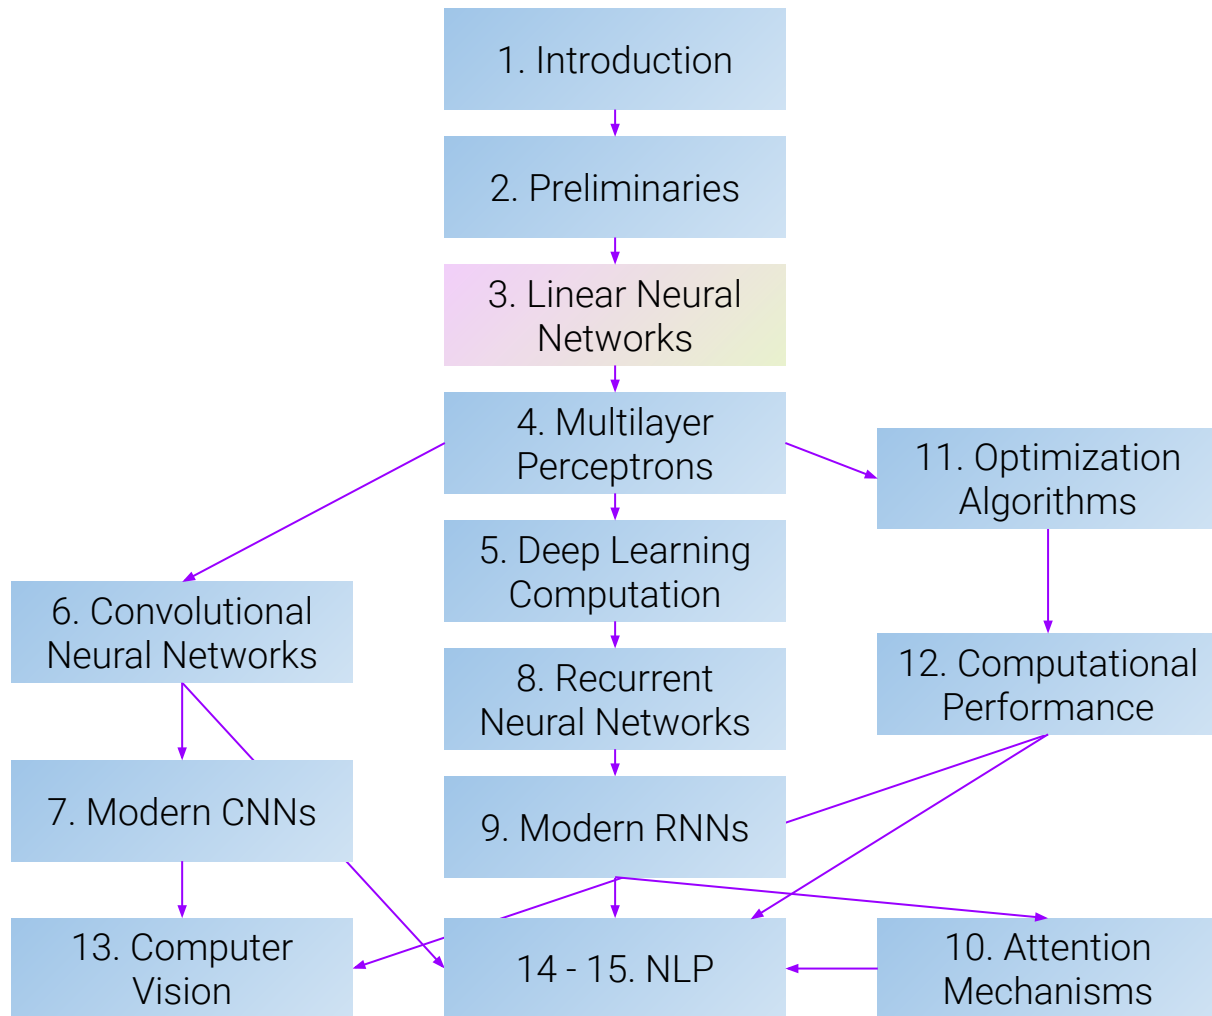


Linear Neural Networks

Session #3

A study group by dair.ai

Chapter 3



Agenda

- Linear Regression
 - Implementation from scratch
 - Concise implementation
- Softmax regression
 - The image classification dataset
 - Implementation from scratch
 - Concise implementation



Linear Regression

- Regression is used for modeling relationship between one or more ***independent*** variables and a ***dependent*** variable
- The aim is to train a model that is able to predict a numerical value
 - Predicting prices (homes, stocks, etc.)
 - Predicting length of stay (patients in hospital)
 - Demand forecasting (sales)
- Linear regression differs from ***classification***
 - Classification aims to predict from set of categories (cat vs. dogs, positive vs. negative)



Basic Elements of Linear Regression

- Assumption:
 - *Linear relationship* between independent variables \mathbf{x} and the dependent variable \mathbf{y}
 - \mathbf{y} can be expressed as a *weighted sum* of the elements in \mathbf{x} , given noise on the observations
 - Noise well behaved (follow Gaussian distribution)
- Example:
 - We would like to estimate price of houses based on **area** and **age**
 - We want to develop a predictive model for predicting house prices
- We need:
 - A **training dataset** or **training set**
 - Rows referred to as *examples*, *data points*, *data instances*, *sample*
 - Dependent variable is called the *label* or *target*
 - Independent variables are called the *features* or *covariates*

area	age	price
15	10	15000
25	15	25000



Notations

- \mathbf{n} denotes the # of examples
- Index the data examples $\mathbf{x}^{(i)} = [x_1^{(i)}, x_2^{(i)}]^\top$ *superscript*
- Corresponding labels $y^{(i)}$



Linear Model

Based on the **linearity assumption** we say that the target is a weighted sum of the features and translation (bias)

$$\text{price} = w_{\text{area}} \cdot \text{area} + w_{\text{age}} \cdot \text{age} + b$$

Weights - Influence of features on the predictions

Bias - Take on this value when features take value of 0

Affine transformation of input features



Linear Model

- **Goal:** choose weights and bias such that on average, predictions of the model fit the true prices observed in the data
- Linear models rely on the ***affine transformation*** specified by the chosen ***weights*** and ***bias***
- So, generally speaking we have $\hat{y} = w_1x_1 + \dots + w_dx_d + b$
- Compact form using vectors: $\hat{y} = \mathbf{w}^\top \mathbf{x} + b$



Linear Model

We refer to features of the entire dataset as:

$$\hat{\mathbf{y}} = \mathbf{X}\mathbf{w} + b$$

Features of dataset

Goal is to find **w** and **b** such that for a new example (from same distribution) and it's label (in expectation) the model makes prediction that produces the lowest error



Linear Models

- We have a formulation based on the linearity assumption
- However,
 - we need a **quality measure** (measure goodness/badness of the model)
 - and a **procedure** to update and improve the model quality

We need a loss function and an optimizer




Loss function

- A function that quantifies the difference between real and predicted value of the target
- The smaller the value of loss the better
- A popular loss function: **squared error**
- Empirical error is a function of the parameters

$$l^{(i)}(\mathbf{w}, b) = \frac{1}{2} \left(\hat{y}^{(i)} - y^{(i)} \right)^2$$

Term cancels
when we take
the derivative
of the loss



$$L(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n l^{(i)}(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (\underbrace{\mathbf{w}^\top \mathbf{x}^{(i)} + b}_{\hat{y}^{(i)}} - y^{(i)})^2$$

Number of examples

When training the model, we seek parameters that minimize the total loss across all training examples

Note the superscript suggests an operation applied to a single example

Average loss on the entire dataset



Loss function

When training the model we are looking for weights and bias (parameters) that ***minimize the total loss on all training examples***

$$\mathbf{w}^*, b^* = \arg \min_{\mathbf{w}, b} L(\mathbf{w}, b)$$



Optimization

- We seek to iteratively reduce the error of the model and improve its quality:
 - Updating the parameters in the direction that incrementally lowers the loss function
- We use the **gradient descent algorithm** to achieve it
- We can directly take the derivative of the average loss on the entire dataset
 - This requires pass over the entire dataset before making a single update
- A better solution is called **minibatch stochastic gradient descent**
 - Sampling random minibatch of examples and
 - Take the derivative of the average loss on the minibatch with regard to the model parameters and then compute the update



Minibatch Stochastic Gradient Descent

$$(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{(\mathbf{w}, b)} l^{(i)}(\mathbf{w}, b)$$

Term multiplied to the gradient (learning rate)

Partial derivative of the average loss of the minibatch

Subtract the result from the current parameters

Minibatch size



Optimization Algorithm

- Initialize model parameters (typically random)
- Sample random minibatches
- Update parameters in the direction of the negative gradient:

$$\begin{aligned}\mathbf{w} &\leftarrow \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{\mathbf{w}} l^{(i)}(\mathbf{w}, b) = \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} \left(\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right), \\ b &\leftarrow b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_b l^{(i)}(\mathbf{w}, b) = b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \left(\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right).\end{aligned}$$



Prediction using Linear Regression Model

- We adjust hyperparameters (e.g., learning rate) assessed on validation set
- We aim to find parameters that achieve low loss on **unseen data**
 - Also referred to **generalization** (discussed later on)
- Given those learned parameters it is now possible to estimate targets given features of a new instance
- We use the squares loss below to quantify goodness/badness of the model
 - Using maximum likelihood estimate principles we get **negative log likelihood**

$$-\log P(\mathbf{y} \mid \mathbf{X}) = \sum_{i=1}^n \frac{1}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)} - b)^2$$



Normal Distribution and Squared loss

Linear regression with the square loss can be motivated by assuming that the observations arise from noisy distributions. Hence,

$$y = \mathbf{w}^\top \mathbf{x} + b + \epsilon \text{ where } \epsilon \sim \mathcal{N}(0, \sigma^2)$$

$$P(y \mid \mathbf{x}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2} (y - \mathbf{w}^\top \mathbf{x} - b)^2\right)$$

$$P(\mathbf{y} \mid \mathbf{X}) = \prod_{i=1}^n p(y^{(i)} \mid \mathbf{x}^{(i)})$$

$$-\log P(\mathbf{y} \mid \mathbf{X}) = \sum_{i=1}^n \frac{1}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)} - b)^2$$



Linear regression as single-layer neural network





Figure source: d2l.ai




Classification

- Classification aims to predict from set of categories (e.g., cat vs. dogs, positive vs. negative)
- Hard assignment of examples to categories (classes)
- Soft assignments; assess probabilities (discussed later on)
- We want a model that estimates the ***conditional probabilities*** with all possible classes
 - Model with multiple outputs (one per class)
- Goal: optimize our parameters to produce probabilities that maximize the likelihood of the observed data
 - One main approach is softmax regression



Softmax Regression

- Supports classification (binary and multiclass)
- Applies softmax function produces outputs as probabilities
- Take logits (output of the linear model) and transform to output probabilities that meet a desired criteria
 - Nonnegativity
 - Differentiable
 - Probabilities sum up to 1



Softmax regression is a single-layer neural network.

Figure source: d2l.ai



Cross-entropy loss

- Used to measure quality of predicted probabilities
- Common loss function used in classification problems
- Computes the expected value of the loss for a distribution over labels
- Concretely, cross-entropy objective
 - Maximizes the likelihood of the observed data
 - Measures difference between two probability distributions
 - Minimize the surprisal require to communicate the labels (refer to information theory)

$$l(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{j=1}^q y_j \log \hat{y}_j$$



Assignment

- Implement backward function
- Calculate the derivatives with respect to the parameters to obtain gradients from scratch and do the update them manually
- Individual assignment
- Solution provided next week



Demo